

Lab 21 - Regularized Regression

Chad Murchison, Guy Twa, Eddie-Williams Owiredo

2023-11-03

- Today's Lab
- Loading Packages and Data
- Getting to Know Your Data
- Fitting regression models
- Comparing model coefficients
- Parameter sweeps using cross-validation
 - LASSO cross validation
- Comparing OLS, LASSO, and Ridge
 - Model coefficients
 - Prediction mean squared error (MSE)
- Conclusions

Today's Lab

Today's lab will introduce you to two supervised machine learning algorithm for regression, Ridge regression (L2) and Lasso (L1) regression. We will learn how to fit regularized regression models and compare them to simple ordinary least squares (OLS) regression. The goal of these regression methods is to perform prediction of outcome variables.

Loading Packages and Data

For today's lab, we will be using a new packages called `glmnet`, which you should go ahead and install prior to loading it in.

```
install.packages("glmnet")
```

In addition to loading our library, data, and setting our working directory, we'll use `set.seed` for replicability.

```
library(glmnet)
set.seed(112233)

setwd("~/Documents/GRD770/Lab 21 - Regularized Regression")

#Example data from glmnet
data(QuickStartExample)
```

Getting to Know Your Data

The dataset we'll be using for this lab is "QuickStartExample" from the glmnet package. This is a randomly generated data set with 100 observations of 20 numerical variables with one numerical outcome variable.

We can get a quick summary of the variables using the `summary()` function:

```
summary(QuickStartExample$x)
```

```
##           V1           V2           V3           V4
## Min.      : -1.8961  Min.      : -2.878895  Min.      : -2.63061  Min.      : -2.05928
## 1st Qu.: -0.4737  1st Qu.: -0.656813  1st Qu.: -0.63942  1st Qu.: -0.61508
## Median :  0.2010  Median : -0.099319  Median :  0.07550  Median :  0.06699
## Mean      :  0.2063  Mean      : -0.005649  Mean      :  0.07555  Mean      :  0.11810
## 3rd Qu.:  0.9952  3rd Qu.:  0.731664  3rd Qu.:  0.85611  3rd Qu.:  0.86831
## Max.      :  3.1619  Max.      :  2.228786  Max.      :  2.55155  Max.      :  3.11291
##           V5           V6           V7           V8
## Min.      : -2.68543  Min.      : -2.78803  Min.      : -2.50549  Min.      : -3.6367
## 1st Qu.: -0.80457  1st Qu.: -0.78204  1st Qu.: -0.66256  1st Qu.: -0.8617
## Median : -0.09333  Median :  0.06646  Median :  0.09519  Median : -0.1421
## Mean      : -0.06115  Mean      :  0.01534  Mean      :  0.03528  Mean      : -0.1218
## 3rd Qu.:  0.58408  3rd Qu.:  0.70799  3rd Qu.:  0.63051  3rd Qu.:  0.5938
## Max.      :  1.81981  Max.      :  2.27527  Max.      :  2.18573  Max.      :  2.6536
##           V9           V10          V11          V12
## Min.      : -2.88985  Min.      : -2.27895  Min.      : -2.2393  Min.      : -2.44167
## 1st Qu.: -0.73077  1st Qu.: -0.54614  1st Qu.: -0.9854  1st Qu.: -0.88268
## Median : -0.09803  Median :  0.02715  Median : -0.3196  Median :  0.03550
## Mean      : -0.08380  Mean      :  0.10124  Mean      : -0.2116  Mean      : -0.02109
## 3rd Qu.:  0.47161  3rd Qu.:  0.81758  3rd Qu.:  0.4524  3rd Qu.:  0.68328
## Max.      :  2.41487  Max.      :  2.68135  Max.      :  2.2831  Max.      :  2.34162
##           V13          V14          V15          V16
## Min.      : -2.71954  Min.      : -2.4622  Min.      : -1.96111  Min.      : -2.22929
## 1st Qu.: -0.65109  1st Qu.: -0.9246  1st Qu.: -0.70695  1st Qu.: -0.60383
## Median : -0.02552  Median : -0.1244  Median : -0.25940  Median : -0.00324
## Mean      : -0.07845  Mean      : -0.1358  Mean      : -0.09221  Mean      :  0.06317
## 3rd Qu.:  0.40112  3rd Qu.:  0.6319  3rd Qu.:  0.53198  3rd Qu.:  0.73559
## Max.      :  2.75606  Max.      :  2.3301  Max.      :  2.25018  Max.      :  2.13515
##           V17          V18          V19          V20
## Min.      : -2.01322  Min.      : -3.4325  Min.      : -3.1944  Min.      : -2.34578
## 1st Qu.: -0.63558  1st Qu.: -0.8642  1st Qu.: -0.5149  1st Qu.: -0.71932
## Median :  0.02783  Median : -0.1694  Median :  0.2007  Median : -0.07920
## Mean      :  0.00815  Mean      : -0.2724  Mean      :  0.1156  Mean      : -0.05225
## 3rd Qu.:  0.58884  3rd Qu.:  0.4617  3rd Qu.:  0.9612  3rd Qu.:  0.50403
## Max.      :  2.56317  Max.      :  2.0789  Max.      :  2.5947  Max.      :  2.51791
```

```
summary(QuickStartExample$y)
```

```
##          V1
## Min.     :-6.8575
## 1st Qu.: -1.6814
## Median :  0.7697
## Mean    :  0.6608
## 3rd Qu.:  2.9038
## Max.     :  6.2659
```

While this data set is pretty abstract, for our purposes great practice data set to show off some of the key concepts with regularized regression. Before we start working, lets coalesce our variables to a single data frame that we'll work with.

```
x <- QuickStartExample$x
y <- QuickStartExample$y

dat <- data.frame(x, Y = y)
dim(dat); head(dat)
```

```
## [1] 100  21
```

```
##          X1          X2          X3          X4          X5          X6          X7
## 1  0.2738562 -0.0366722  0.8547269  0.9675242  1.4154898  0.5234059  0.5626882
## 2  2.2448169 -0.5460300  0.2340651 -1.3350304  1.3130758  0.5212746 -0.6100346
## 3 -0.1254230 -0.6068782 -0.8539217 -0.1487772 -0.6646828  0.6066164  0.1617207
## 4 -0.5435734  1.1083583 -0.1042480  1.0165262  0.6999042  1.6550164  0.4899635
## 5 -1.4593984 -0.2744945  0.1119060 -0.8517877  0.3152839  1.0507493  1.3863575
## 6  1.0632081 -0.7535232 -1.3825534  1.0762270  0.3700331  1.4987212 -0.3604525
##          X8          X9          X10          X11          X12          X13
## 1  1.11122333  1.6408214  0.6187067  0.99993483 -0.07841916 -0.60332610
## 2 -0.86139651 -0.2704635  0.2300825 -0.10570898  0.16314122  0.76207661
## 3 -0.86272165  0.6042102  1.1939768  0.50125094 -0.94520592  0.39890263
## 4  0.02338209  0.2560304 -0.1273140 -0.06262846  0.64195468  0.07548198
## 5  0.28450104  1.1404976  2.6813460 -1.01473386  0.36704372  1.73759745
## 6 -0.21421571  1.8266483  0.8711225 -0.06372928  1.00507110  0.31520786
##          X14          X15          X16          X17          X18          X19
## 1  0.03323168 -0.7008845  1.1578379  1.4578156  0.77490699 -1.2685177
## 2  0.67812003 -0.5282670 -0.8791548 -0.4729134 -1.11717309 -0.7377321
## 3 -0.76478505  1.2854019  0.6448767  0.1792455  0.04473756  1.1053071
## 4 -1.37846440 -1.0247390 -2.1183615 -0.4695345  0.69779616  0.8656362
## 5 -1.26612706  1.4519995 -0.7894756 -0.9843860 -1.63700993 -0.5829168
## 6  0.53671846 -1.2624227 -1.5848662 -0.6314991 -1.87874350  0.4504287
##          X20          Y
## 1  1.9935800 -1.2748860
## 2 -1.0787929  1.8434251
## 3  0.3040545  0.4592363
## 4 -0.7894895  0.5640407
## 5 -1.5289104  1.8729633
## 6  1.4422643  0.5275317
```

Fitting regression models

We will use the `glmnet()` function from the `glmnet` package to fit our regularized regression models. This function has a lot of different parameters but we'll focus on the following:

- `x` = predictor variable data, as a matrix or dataframe
- `y` = response variable data, as a vector
- `family` = the model family, use "gaussian" for a continuous outcome variable (the OLS equivalent)
- `alpha` = the mixing parameter for the elastic net, 0 for Ridge, 1 for LASSO
- `lambda` = vector of lambda shrinkage penalty parameters
- `nlambda` = the number of shrinkage penalty parameters we want to search over
- `standardize` = whether or not to standardize the data before fitting the model

The mixing parameter, `alpha` determines the relative weight of the ridge and LASSO regression penalties. We'll be using values of either 0 for ridge regression or 1 for LASSO regression. When you use a mixture of these its called an elastic net. Today though, we'll focus on getting a sense of the differences between ridge and LASSO.

Lambda is our shrinkage penalty applied to our coefficients. Remember, if `lambda` is too large, all coefficients shrink to zero while setting `lambda` to zero leads to OLS regression. If we want to make our own grid sweep of `lambda` values we can create a vector and supply it to the `lambda` parameter. That can then be passed to the `lambda` argument. Alternatively we can let `glmnet()` come up with a set of `lambda` values, in which case we'll need the `nlambda` parameter to set the size for the range of `lambda` values to try.

Standardizing our inputs in advance is important as these methods are not scale equivalent like OLS regression. This is especially important when input variables are measured on different scales (e.g. height and weight). Thus, we set `standardize` to `TRUE`.

All of the other parameters are about particular operations of the function and beyond our scope here.

Let's fit a Ridge (L2) regression model by setting `alpha` as 0, and a LASSO (L1) regression model by setting `alpha` as 1. (The parameters are listed out for our Ridge regression just for illustration. These are all the default values, so they are the same when we fit the LASSO as well.)

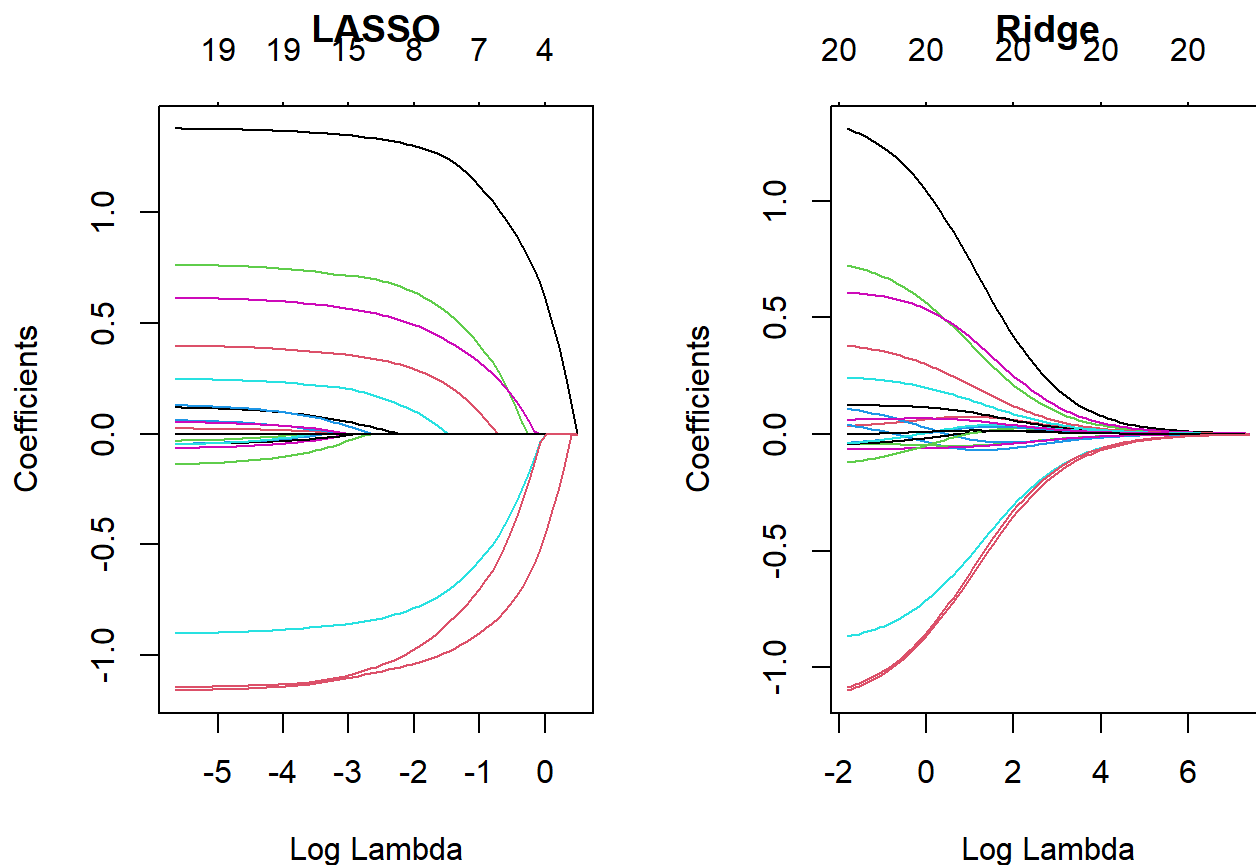
```
#Fit a ridge regression
fit_l2 <- glmnet(x, y,
                family = "gaussian",
                alpha = 0,
                nlambda = 100,
                lambda = NULL,
                standardize = TRUE)

#Fit a LASSO by setting alpha to 1
fit_l1 <- glmnet(x, y, alpha = 1)
```

Comparing model coefficients

We can compare the effects of the different regularization penalties by plotting the predictor variable coefficient values against `lambda`.

```
par(mfrow = c(1,2))
plot(fit_l1, xvar = "lambda", main = "LASSO")
plot(fit_l2, xvar = "lambda", main = "Ridge")
```



We see how as lambda increases, the L2 norm of the LASSO shrinks coefficient to 0 while the Ridge only approaches 0. Meanwhile, when the lines are flat early on, this is essentially the OLS coefficient values. An important point about both plots the number of coefficients labeled up top. LASSO regression drops coefficients as they reach 0, but Ridge regression always retains coefficients because they never reach 0.

We can also look at these results over the iterations of lambda values.

```
print(fit_l1); print(fit_l2)
```

```
##
## Call:  glmnet(x = x, y = y, alpha = 1)
##
##      Df  %Dev  Lambda
## 1    0  0.00 1.63100
## 2    2  5.53 1.48600
## 3    2 14.59 1.35400
## 4    2 22.11 1.23400
## 5    2 28.36 1.12400
## 6    2 33.54 1.02400
## 7    4 39.04 0.93320
## 8    5 45.60 0.85030
## 9    5 51.54 0.77470
## 10   6 57.35 0.70590
## 11   6 62.55 0.64320
## 12   6 66.87 0.58610
## 13   6 70.46 0.53400
## 14   6 73.44 0.48660
## 15   7 76.21 0.44330
## 16   7 78.57 0.40400
## 17   7 80.53 0.36810
## 18   7 82.15 0.33540
## 19   7 83.50 0.30560
## 20   7 84.62 0.27840
## 21   7 85.55 0.25370
## 22   7 86.33 0.23120
## 23   8 87.06 0.21060
## 24   8 87.69 0.19190
## 25   8 88.21 0.17490
## 26   8 88.65 0.15930
## 27   8 89.01 0.14520
## 28   8 89.31 0.13230
## 29   8 89.56 0.12050
## 30   8 89.76 0.10980
## 31   9 89.94 0.10010
## 32   9 90.10 0.09117
## 33   9 90.23 0.08307
## 34   9 90.34 0.07569
## 35  10 90.43 0.06897
## 36  11 90.53 0.06284
## 37  11 90.62 0.05726
## 38  12 90.70 0.05217
## 39  15 90.78 0.04754
## 40  16 90.86 0.04331
## 41  16 90.93 0.03947
## 42  16 90.98 0.03596
## 43  17 91.03 0.03277
## 44  17 91.07 0.02985
## 45  18 91.11 0.02720
## 46  18 91.14 0.02479
## 47  19 91.17 0.02258
## 48  19 91.20 0.02058
```

```
## 49 19 91.22 0.01875
## 50 19 91.24 0.01708
## 51 19 91.25 0.01557
## 52 19 91.26 0.01418
## 53 19 91.27 0.01292
## 54 19 91.28 0.01178
## 55 19 91.29 0.01073
## 56 19 91.29 0.00978
## 57 19 91.30 0.00891
## 58 19 91.30 0.00812
## 59 19 91.31 0.00739
## 60 19 91.31 0.00674
## 61 19 91.31 0.00614
## 62 20 91.31 0.00559
## 63 20 91.31 0.00510
## 64 20 91.31 0.00464
## 65 20 91.32 0.00423
## 66 20 91.32 0.00386
## 67 20 91.32 0.00351
```

```
##
## Call:  glmnet(x = x, y = y, family = "gaussian", alpha = 0, nlambda = 100,      lambd
a = NULL, standardize = TRUE)
##
##      Df  %Dev  Lambda
## 1    20  0.00 1631.00
## 2    20  0.45 1486.00
## 3    20  0.49 1354.00
## 4    20  0.54 1234.00
## 5    20  0.59 1124.00
## 6    20  0.65 1024.00
## 7    20  0.71  933.20
## 8    20  0.78  850.30
## 9    20  0.86  774.70
## 10   20  0.94  705.90
## 11   20  1.03  643.20
## 12   20  1.13  586.10
## 13   20  1.24  534.00
## 14   20  1.36  486.60
## 15   20  1.49  443.30
## 16   20  1.63  404.00
## 17   20  1.79  368.10
## 18   20  1.96  335.40
## 19   20  2.14  305.60
## 20   20  2.35  278.40
## 21   20  2.57  253.70
## 22   20  2.82  231.20
## 23   20  3.09  210.60
## 24   20  3.38  191.90
## 25   20  3.70  174.90
## 26   20  4.04  159.30
## 27   20  4.42  145.20
## 28   20  4.84  132.30
## 29   20  5.29  120.50
## 30   20  5.77  109.80
## 31   20  6.31  100.10
## 32   20  6.88   91.17
## 33   20  7.51   83.07
## 34   20  8.19   75.69
## 35   20  8.93   68.97
## 36   20  9.72   62.84
## 37   20 10.58   57.26
## 38   20 11.51   52.17
## 39   20 12.50   47.54
## 40   20 13.58   43.31
## 41   20 14.73   39.47
## 42   20 15.96   35.96
## 43   20 17.27   32.77
## 44   20 18.68   29.85
## 45   20 20.17   27.20
## 46   20 21.75   24.79
## 47   20 23.42   22.58
```


##	48	20 25.19	20.58
##	49	20 27.04	18.75
##	50	20 28.98	17.08
##	51	20 31.01	15.57
##	52	20 33.12	14.18
##	53	20 35.30	12.92
##	54	20 37.55	11.78
##	55	20 39.86	10.73
##	56	20 42.22	9.78
##	57	20 44.62	8.91
##	58	20 47.05	8.12
##	59	20 49.50	7.39
##	60	20 51.95	6.74
##	61	20 54.39	6.14
##	62	20 56.81	5.59
##	63	20 59.19	5.10
##	64	20 61.52	4.64
##	65	20 63.78	4.23
##	66	20 65.98	3.86
##	67	20 68.09	3.51
##	68	20 70.11	3.20
##	69	20 72.04	2.92
##	70	20 73.85	2.66
##	71	20 75.56	2.42
##	72	20 77.16	2.21
##	73	20 78.64	2.01
##	74	20 80.01	1.83
##	75	20 81.27	1.67
##	76	20 82.42	1.52
##	77	20 83.47	1.39
##	78	20 84.42	1.26
##	79	20 85.27	1.15
##	80	20 86.04	1.05
##	81	20 86.72	0.96
##	82	20 87.32	0.87
##	83	20 87.86	0.79
##	84	20 88.33	0.72
##	85	20 88.75	0.66
##	86	20 89.11	0.60
##	87	20 89.43	0.55
##	88	20 89.70	0.50
##	89	20 89.94	0.45
##	90	20 90.14	0.41
##	91	20 90.32	0.38
##	92	20 90.47	0.34
##	93	20 90.60	0.31
##	94	20 90.71	0.28
##	95	20 90.81	0.26
##	96	20 90.89	0.24
##	97	20 90.96	0.22
##	98	20 91.01	0.20

```
## 99 20 91.06 0.18
## 100 20 91.10 0.16
```

We see the L2 converged more quickly while the L1 probably need more lambdas to explain the same amount of variance.

We can extract the coefficients at a specific value of lambda defined as `s` using the `coef()` function

```
coef(fit_l1, s = 0.1)
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 0.150928072
## V1          1.320597195
## V2          .
## V3          0.675110234
## V4          .
## V5         -0.817411518
## V6          0.521436671
## V7          0.004829335
## V8          0.319415917
## V9          .
## V10         .
## V11         0.142498519
## V12         .
## V13         .
## V14        -1.059978702
## V15         .
## V16         .
## V17         .
## V18         .
## V19         .
## V20        -1.021873704
```

```
coef(fit_l2, s = 0.1)
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)  0.14576689
## V1          1.30922437
## V2          0.03496846
## V3          0.72391240
## V4          0.03882705
## V5         -0.86710569
## V6          0.60697109
## V7          0.12355737
## V8          0.37889309
## V9         -0.03973640
## V10         0.10841981
## V11         0.24189927
## V12        -0.06661643
## V13        -0.04268166
## V14        -1.09804121
## V15        -0.12176667
## V16        -0.03711366
## V17        -0.04019624
## V18         0.06146105
## V19        -0.00179925
## V20        -1.08563245
```

Again we see how the L2 will completely shrink (and select) coefficients

Parameter sweeps using cross-validation

If we want to find an ideal lambda across a range, we can use a built in cross-validation command. To do this, we're going to want to build a test set and training, we'll use 60% of the QuickData dataframe.

We can select our training samples by randomly selecting 60 samples, and designating the remaining observations as testing samples.

```
train_vec <- sort(sample(1:nrow(dat), floor(nrow(dat) * 0.60)))
test_vec  <- setdiff(1:nrow(dat), train_vec)
```

We'll also make our own vector of lambda values to sweep.

```
lambda_sweep <- 10 ^ seq(5, -5, length = 500)
```

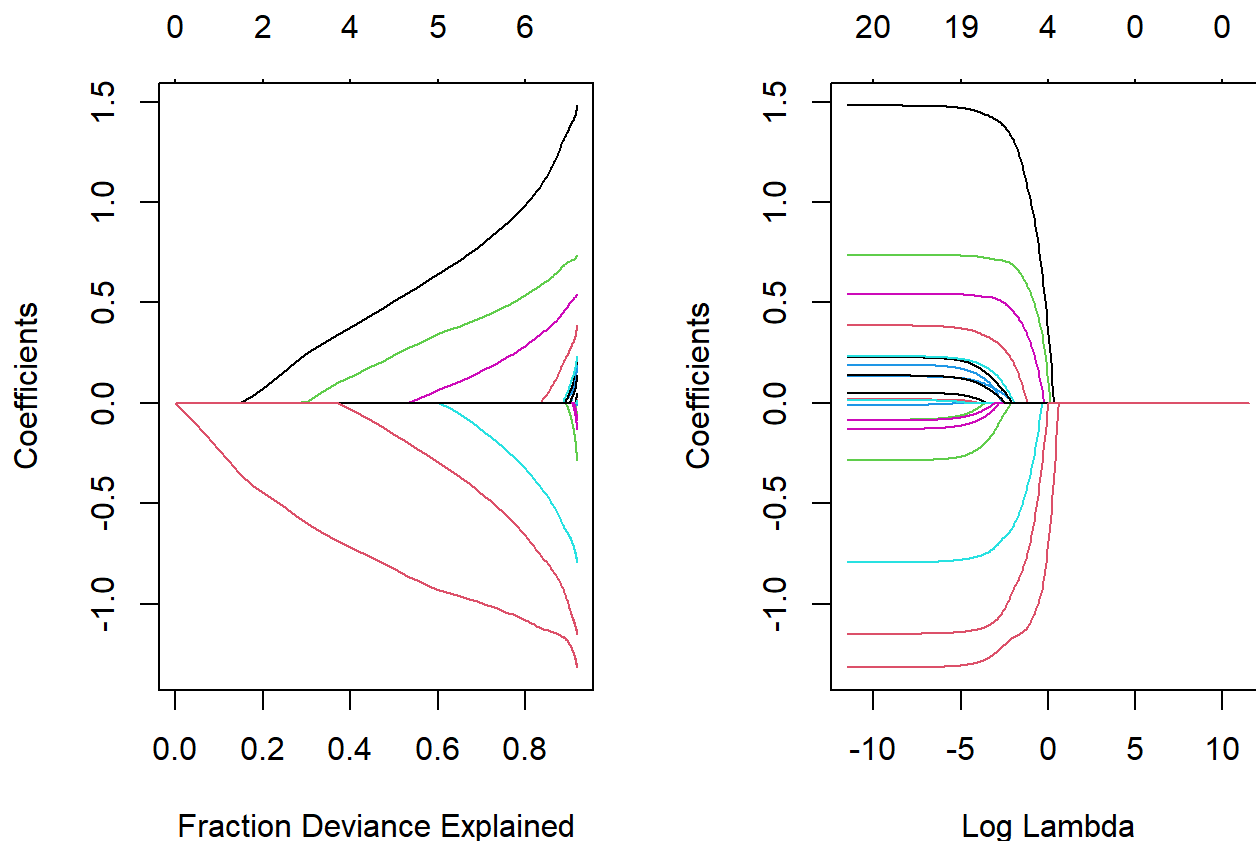
The lambda sweep vector effectively gives a set of lambda's ranging from intercept only ($\lambda = 10^5$) to basic OLS ($\lambda = 0.00001$).

LASSO cross validation

Let's start with the full set using LASSO for ease of display, and plot the coefficients against the explained deviance as well as the range of lambda values.

```
# fit LASSO regression using the training
fit_l1_sweep <- glmnet(x[train_vec,],
                      y[train_vec],
                      alpha = 1,
                      lambda = lambda_sweep)

par(mfrow = c(1,2))
plot(fit_l1_sweep, xvar = "dev", label = TRUE)
plot(fit_l1_sweep, xvar = "lambda")
```




From the left plot, it appears we only need about 7 predictors to explain over 90% of the deviance. From our right plot, our sweep of lambda's mainly shows when lambda essentially becomes too large. But what's the most ideal tuning?

For that we can test with cross validation using the `cv.glmnet()` function. The most important parameter here is `nfolds`, where we set the number of folds for cross validation. For now, we'll leave that at the default 10.

```
fit_l1_cv <- cv.glmnet(x[train_vec,], y[train_vec], alpha = 1)
```

We can plot the MSE as a function of $\log(\lambda)$ to start to see what the sweet spot is going to be where we have a minimum MSE with the fewest variables. Again, it's a tradeoff of bias and variance

```
plot(fit_l1_cv)
```



We see the best lambda value is going to be around 0.1008 - 0.2555 at $\log(\lambda)$ of -2.294617 to -1.364533 and corresponds to roughly 9-11 predictors. We can select the lambda with the lowest MSE directly from the CV object by accessing the `lambda.min` variable.

```
best_l1_lambda <- fit_l1_cv$lambda.min
best_l1_lambda
```

```
## [1] 0.1007692
```

Now that we have our model, we can evaluate how well it makes predictions using our testing data we set aside earlier. To make predictions, we use the `predict()` function. In our `predict()` call we need to specify our model, our lambda value, and our testing data predictor variable data.

```
pred_l1_cv <- predict(fit_l1_sweep,      # model
                      s = best_l1_lambda, # selected lambda value
                      newx = x[test_vec,]) # testing data predictor observations
```

Finally, we can look at the mean squared error (MSE) of our model outcome predictions to the real values.

```
mean((pred_l1_cv - y[test_vec])^2)
```

```
## [1] 1.007394
```

Comparing OLS, LASSO, and Ridge

Now we can compare against the expected OLS output we can see the similarity with small lambda. Let's start by building out a cross-validated dataset for Ridge to go with our Lasso using the same method as before.

Ridge cross validation

```
# Fit models for our range of lambda values
fit_l2_sweep <- glmnet(x[train_vec,], y[train_vec],
                      alpha = 0, lambda = lambda_sweep)

# Select the best lambda using cross-validation
fit_l2_cv <- cv.glmnet(x[train_vec,], y[train_vec], alpha = 0)
best_l2_lambda <- fit_l2_cv$lambda.min

# Make predictions
pred_l2_cv <- predict(fit_l2_sweep, s = best_l2_lambda, newx = x[test_vec,])
```

OLS model

Next, let's build our OLS model. We can build this using the basic `lm()` function. We can use the `predict()` function with this model as well, by supplying our model and testing data.

```
fit_ols <- lm(Y ~ ., data = dat[train_vec,])
pred_ols <- predict(fit_ols, newdata = dat[test_vec,])
```

Model coefficients

Now that we have all three models, let's put their coefficients side by side to compare. Here we use the `cbind()` function to column-wise bind our sets of coefficients together.

```
coef_results <- cbind(OLS = coef(fit_ols),
                     Ridge_L2 = coef(fit_l2_sweep, s = best_l2_lambda)[,1],
                     LASSO_L1 = coef(fit_l1_sweep, s = best_l1_lambda)[,1])
coef_results
```

	OLS	Ridge_L2	LASSO_L1
## (Intercept)	0.05821383	0.102880176	0.17014352
## X1	1.48091879	1.322703316	1.35401874
## X2	0.01709023	0.044406668	0.00000000
## X3	0.73784355	0.703911301	0.70065873
## X4	0.13231888	0.091561381	0.02718769
## X5	-0.79405163	-0.717792595	-0.64943556
## X6	0.54314247	0.531302527	0.48226569
## X7	0.22752049	0.198366755	0.03250319
## X8	0.38581490	0.306514059	0.24401196
## X9	-0.08452980	-0.037763656	0.00000000
## X10	0.19059793	0.098247122	0.00000000
## X11	0.23233986	0.186836195	0.06042086
## X12	-0.13028683	-0.110309384	0.00000000
## X13	0.05012370	0.031315284	0.00000000
## X14	-1.31681810	-1.192941934	-1.18867733
## X15	-0.28633018	-0.197201239	-0.02353272
## X16	-0.01020831	0.025689534	0.00000000
## X17	0.01383254	0.001036487	0.00000000
## X18	-0.08699664	-0.061251733	0.00000000
## X19	0.13672434	0.106198897	0.00000000
## X20	-1.14993373	-1.036652851	-0.99059896

Again, we see how the Ridge will not zero out a coefficient set but LASSO will.

Prediction mean squared error (MSE)

Now, let's compare the prediction MSE for our three models.

```
mean((pred_ols - y[test_vec])^2)
```

```
## [1] 1.053824
```

```
mean((pred_l2_cv - y[test_vec])^2)
```

```
## [1] 0.9452646
```

```
mean((pred_l1_cv - y[test_vec])^2)
```

```
## [1] 1.007394
```

We can see how the MSE is improved for regularized regression and that LASSO and Ridge largely in agreement.

Similarly, when lambda is prohibitively large we approach an intercept-only model

```
fit_ols_intercept <- lm(Y ~ 1, data=dat[train_vec,])

coef_intercept <- cbind(OLS = coef(fit_ols_intercept),
                        Ridge_L2 = coef(fit_l2_sweep, s = 10^5),
                        Ridge_L1 = coef(fit_l1_sweep, s = 10^5))

coef_intercept
```

```
## 21 x 3 sparse Matrix of class "dgCMatrix"
##              OLS              s1              s1
## (Intercept) 0.3632427 3.632368e-01 0.3632427
## V1          0.3632427 4.169325e-05 .
## V2          0.3632427 8.956484e-06 .
## V3          0.3632427 3.339678e-05 .
## V4          0.3632427 -7.233409e-06 .
## V5          0.3632427 -2.923502e-05 .
## V6          0.3632427 2.851030e-05 .
## V7          0.3632427 2.734383e-06 .
## V8          0.3632427 4.895044e-06 .
## V9          0.3632427 4.115155e-06 .
## V10         0.3632427 -1.121531e-05 .
## V11         0.3632427 2.956202e-06 .
## V12         0.3632427 -6.091761e-06 .
## V13         0.3632427 -3.564070e-06 .
## V14         0.3632427 -4.777718e-05 .
## V15         0.3632427 6.817468e-06 .
## V16         0.3632427 6.509293e-06 .
## V17         0.3632427 -4.473759e-06 .
## V18         0.3632427 1.454444e-06 .
## V19         0.3632427 3.531605e-06 .
## V20         0.3632427 -3.538384e-05 .
```

We can also see the predictions are just terrible

```
pred_ols_intercept <- predict(fit_ols_intercept, newdata = dat[test_vec,])
mean((pred_ols_intercept - y[test_vec])^2)
```

```
## [1] 8.626318
```

At this stage, we should note this dataset in particular would tend to favor OLS if we added more data to the training set. If our sampling fraction were closer to 90% this would almost entirely favor OLS in terms of the MSE. But this is only because we have so many samples to choose from and a fairly informative dataset. This would not be the case with real world datasets or when our data is especially wide.

We can see what happens with wider datasets, by building models with only the first 10 observations. Now we see how OLS completely collapses but both regularized regressions will find values


```

fit_vif_ols <- lm(Y ~ ., dat[1:10,])
fit_vif_l1 <- glmnet(x[1:10,], y[1:10], alpha=0)
fit_vif_l2 <- glmnet(x[1:10,], y[1:10], alpha=1)

coef_wide <- cbind(OLS = coef(fit_vif_ols),
                  Ridge_L2 = coef(fit_vif_l2, s=0.1),
                  LASSO_L1 = coef(fit_vif_l1, s=0.1))

coef_wide

```

```

## 21 x 3 sparse Matrix of class "dgCMatrix"
##              OLS              s1              s1
## (Intercept)  15.833565  0.87232540  0.8530542306
## V1          -159.824673  .          -0.0031086223
## V2           20.364433  .          -0.0162474069
## V3          -64.113457  0.04836284  0.0595702368
## V4          -4.507769 -0.50353361 -0.0626598457
## V5          188.124659 -0.14342495 -0.0759687201
## V6          -119.552068  .          -0.0416955421
## V7          -150.862726 -0.20354188 -0.0604634680
## V8          -122.147841  .          0.0005305579
## V9           63.063623  .          -0.0019866779
## V10           NA      .          0.0117105106
## V11           NA -0.41377572 -0.0936635964
## V12           NA  0.48589410  0.0918983624
## V13           NA  .          0.0120905264
## V14           NA  .          -0.0498159451
## V15           NA  .          -0.0115419195
## V16           NA  .          -0.0404624536
## V17           NA  .          -0.0552171118
## V18           NA  .          -0.0174065814
## V19           NA  0.13223663  0.0420557018
## V20           NA -0.19288654 -0.0570310321

```

Conclusions

Now we've got a sense for the differences between regularized regression methods and OLS regression. Ridge (L2) regression will shrink coefficient estimates towards zero, but will always retain every variable. LASSO regression (L1) will shrink estimates all the way to zero, effectively removing them from the model. Remember: the goal for these models is not **interpretation** of predictor and outcome variables as with OLS, but **prediction** of outcome variables from new data. Regularized regression methods produce biased coefficient estimates in service of better prediction.